

Figure 1 is a scatter plot with error bars showing the median percentile rank of the word with the highest number of correct associations as a function of gallery size N (words). The y-axis is labeled 'Median percentile rank' and ranges from 0.0 to 0.5. The x-axis is labeled 'Gallery size N (words)' and has values 200, 500, 1000, 2000, and 5000. A dashed horizontal line at 0.5 represents the chance level. Data points are shown for $N=200, 500, 1000, 2000, 5000$, with error bars and significance markers (***) indicating $p < 0.001$. The median rank decreases as N increases, approaching 0.0.

Figure 1 is a line plot showing Top-k retrieval accuracy versus Rank k (gallery N=5000). The y-axis is logarithmic, ranging from 10⁻³ to 10⁻¹. The x-axis shows Rank k values: 1, 5, 10, 50, and 100. A solid black line with error bars represents the proposed method, which shows significantly higher accuracy than the chance baseline (dashed grey line). The proposed method's accuracy increases from approximately 0.03 at k=1 to nearly 0.4 at k=100. The chance baseline starts at approximately 0.0002 at k=1 and reaches approximately 0.02 at k=100. Significance markers (***) indicate statistical significance for the proposed method's performance at k=1, 5, 10, 50, and 100.

Figure 1 is a scatter plot with error bars showing the median percentile rank of word frequency for two sets: In-vocab and Held-out (zero-shot). The y-axis is labeled 'Median percentile rank' and has a logarithmic scale with major ticks at 0.01, 0.1, and 0.5. A dashed horizontal line is at 0.5. The x-axis has two categories: 'In-vocab' and 'Held-out (zero-shot)'. For each category, there are multiple data points represented by colored circles (blue, green, yellow, orange, purple, grey) and a thick black line connecting the median values. Error bars represent the standard deviation. A bracket with three asterisks (***) indicates a significant difference between the two sets.

Set	Median Percentile Rank (approx.)
In-vocab	0.018
Held-out (zero-shot)	0.035

Figure 10: Difference in word frequency (Wu-Palmer) between matched null and neural retrieval conditions. The plot shows that the difference in word frequency is significantly higher in the neural retrieval condition compared to the matched null condition (***).

Box plot showing nDCG@100 (independent WordNet grade) for Matched null and Neural retrieval methods. The Neural retrieval method shows a significantly higher median nDCG@100 (approximately 0.645) compared to the Matched null method (approximately 0.625). A thick black line connects the medians, and a bracket with *** indicates a significant difference.

Method	Median	Q1	Q3	Min	Max
Matched null	~0.625	~0.621	~0.630	~0.590	~0.642
Neural retrieval	~0.645	~0.638	~0.654	~0.610	~0.674

[illegible]